

# Project Report: Automated Deployment of Scalable Applications on AWS EC2 with Kubernetes and Argo CD

DevOpsEngineer/VaibhavAgarwal

## INTRODUCTION

---

In the modern landscape of software engineering, the shift from monolithic architectures to microservices has necessitated robust orchestration tools and automated deployment strategies. Manual deployment of distributed systems is error-prone and inefficient. This project addresses these challenges by implementing a **GitOps** workflow using **Kubernetes** and **Argo CD**. By treating infrastructure as code and using Git as the single source of truth, I established a deployment pipeline that is both auditable and self-healing.

## ABSTRACT

---

This project demonstrates the end-to-end deployment of a multi-tier "Voting Application" on a Kubernetes cluster hosted on AWS EC2. The application consists of five distinct microservices: a Python-based voting interface, a Redis queue, a .NET worker, a PostgreSQL database, and a Node.js results interface.

The core objective was to move beyond manual 'kubectl' application and implement a **GitOps** continuous delivery model. I utilized **Kind** (Kubernetes in Docker) to bootstrap a lightweight cluster and **Argo CD** to synchronize the cluster state with a GitHub repository. The project successfully achieved automated syncing, where changes pushed to the repository were immediately reflected in the live environment without manual intervention.

## TOOLS USED

---

The following technologies were utilized to build the infrastructure and deployment pipeline:

- **AWS EC2 (Ubuntu 22.04)**: Served as the host infrastructure, simulating a cloud environment.
- **Docker**: Used as the container runtime to host the Kind cluster nodes.
- **Kind (Kubernetes in Docker)**: Used to bootstrap a lightweight, production-like Kubernetes cluster.
- **Kubernetes (K8s)**: The orchestration platform managing the microservices, deployments, and services.
- **Argo CD**: The GitOps continuous delivery tool responsible for monitoring the Git repository and syncing manifests.
- **GitHub**: Hosted the source code and Kubernetes manifest files (YAML), acting as the "Source of Truth."
- **Kubernetes Dashboard**: Provided a web-based user interface for visualizing cluster resources.

## **STEPS INVOLVED IN BUILDING THE PROJECT**

---

### **Step 1: Infrastructure Provisioning**

An AWS EC2 instance (t2.medium) was launched with Ubuntu 22.04. Essential dependencies were installed, including Docker (for container runtime) and 'kubectl' (for cluster interaction).

### **Step 2: Cluster Initialization**

I utilized 'kind' to bootstrap the Kubernetes cluster. A configuration file ('config.yml') was used to define the cluster nodes. Verification was performed using 'kubectl get nodes' to ensure the control plane was ready.

### **Step 3: Kubernetes Dashboard Setup**

To gain visual observability into the cluster, I deployed the standard Kubernetes Dashboard using its official YAML manifest. I created an 'admin-user' ServiceAccount with cluster-wide permissions and generated a secure token for authentication. The dashboard was accessed via port-forwarding, allowing me to monitor workloads, logs, and resource usage directly from the browser.

### **Step 4: Argo CD Installation**

Argo CD was deployed into a dedicated namespace ('argocd') using its official manifest. To make the Argo CD UI accessible from the internet, I patched the service type to 'Node-Port' and utilized port-forwarding techniques to tunnel traffic from the EC2 instance to the browser.

### **Step 5: GitOps Configuration**

I accessed the Argo CD UI and configured a "New App." I connected the application to the GitHub repository hosting the microservices' YAML files ('vote-deployment.yaml', 'redis.yaml', etc.). The sync policy was set to "Automatic" with "Self-Heal" enabled.

### **Step 6: Deployment and Verification**

Upon creation, Argo CD automatically pulled the manifests from GitHub and deployed the five microservices. I verified the deployment by accessing the Voting App and Result App via their respective ports. I further tested the GitOps flow by modifying the replica count in GitHub and observing the automatic scale-up in the cluster.

## **CONCLUSION**

---

This project successfully demonstrated the power of GitOps in managing cloud-native applications. By integrating **Argo CD** with **Kubernetes**, the need for manual deployment scripts was eliminated, reducing the risk of human error.

The implemented architecture ensures that the cluster state always matches the configuration in Git. Features like **Self-Healing** (where Argo CD automatically recreates deleted resources) and **Automated Syncing** significantly improved the reliability and maintainability of the deployment process. This setup provides a solid foundation for scalable, production-grade DevOps pipelines.